

TensorFlow Functions for Image Generation

1. `tf.image.draw_bounding_boxes`

Visualizes predicted bounding boxes on images to debug results fast.

Example: `tf.image.draw_bounding_boxes(img, boxes, colors)`

Why: If you can't visualize boxes, you're flying blind.

2. `tf.image.non_max_suppression`

Removes overlapping duplicate boxes keeping the best predictions.

Example: `tf.image.non_max_suppression(boxes, scores, max_output_size=50)`

Why: Every real detector needs NMS or outputs look like spaghetti.

3. `tf.image.resize`

Resizes input frames before feeding to a model (YOLO, SSD, etc.).

Example: `tf.image.resize(img, (300, 300))`

Why: Detection models demand fixed resolutions.

4. `tf.image.combined_non_max_suppression`

Batch-friendly NMS for modern detectors (multi-class + multi-box).

Example: `tf.image.combined_non_max_suppression(boxes, scores, ...)`

Why: Used in production-grade TF models like EfficientDet.

5. `tf.keras.applications.EfficientNetB0 (as backbone)`

Used as a feature extractor for detectors (EfficientDet).

Example: `backbone = EfficientNetB0(include_top=False)`

Why: Detection needs good feature maps, not just logits.

6. `tf.keras.layers.Conv2D`

Core layer for anchor feature generation + box regressions.

Example: `layers.Conv2D(256, 3, padding="same")`

Why: Detection heads = small conv towers.

7. `tf.keras.layers.UpSampling2D`

Upscales feature maps to build feature pyramids.

Example: `layers.UpSampling2D(size=2)`

Why: Multi-scale detection matters (small objects are annoying).

8. tf.keras.layers.Concatenate

Merges feature maps across scales (FPN/biFPN style).

Example: layers.Concatenate()([p3, p4])

Why: No concatenation → multi-scale detection dies.

9. tf.image.random_brightness (*augmentation*)

Adds brightness noise to strengthen robustness.

Example: tf.image.random_brightness(img, 0.3)

Why: Real-world lighting sucks; augment for it.

10. tf.image.random_flip_left_right (*augmentation*)

Flips objects + boxes together during training.

Example: tf.image.random_flip_left_right(img)

Why: Cheap augmentation that avoids overfitting.